Self-Organizing Data Containers

Thomas Glas

Technische Universität München

TUM School of Computation, Information and Technology

Lehrstuhl für Datenbanksystem

München, 14. März 2023

Motivation

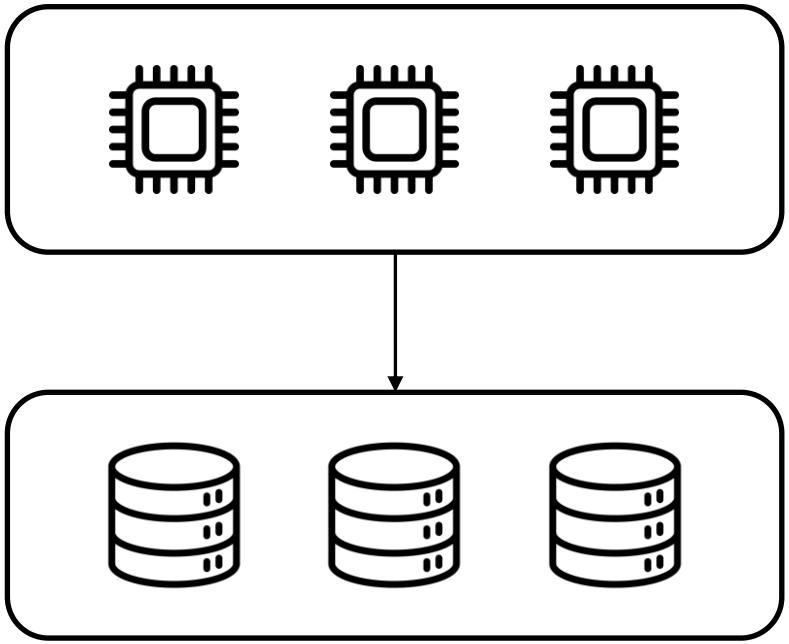

How can we minimize data transfer?

➡ only retrieve data needed for query

How can we create a storage layer with rich metadata to support this?

Self-Organizing Data Containers

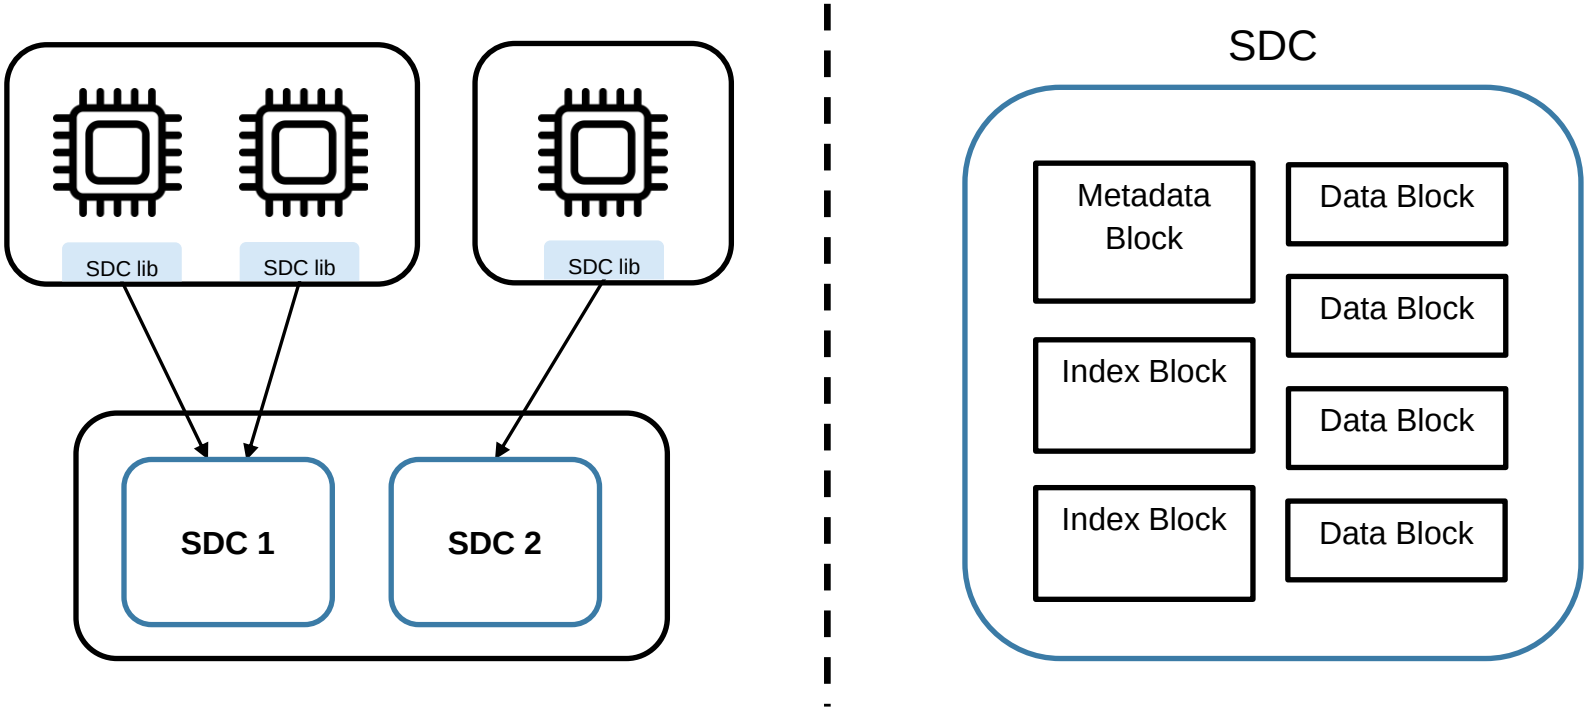

Partitioning Strategies

Column Range Partitioning

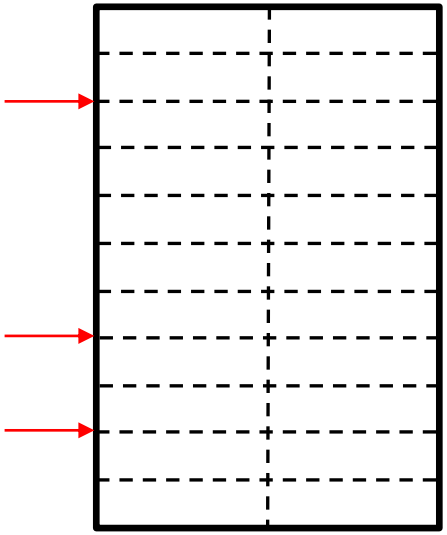

Qd-Tree

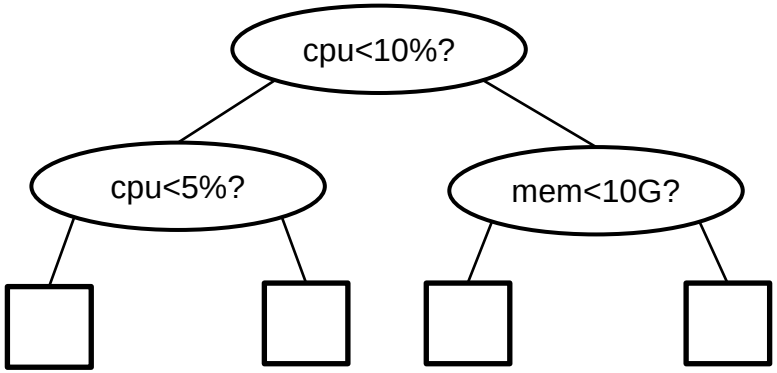

Self-Organizing Data Containers

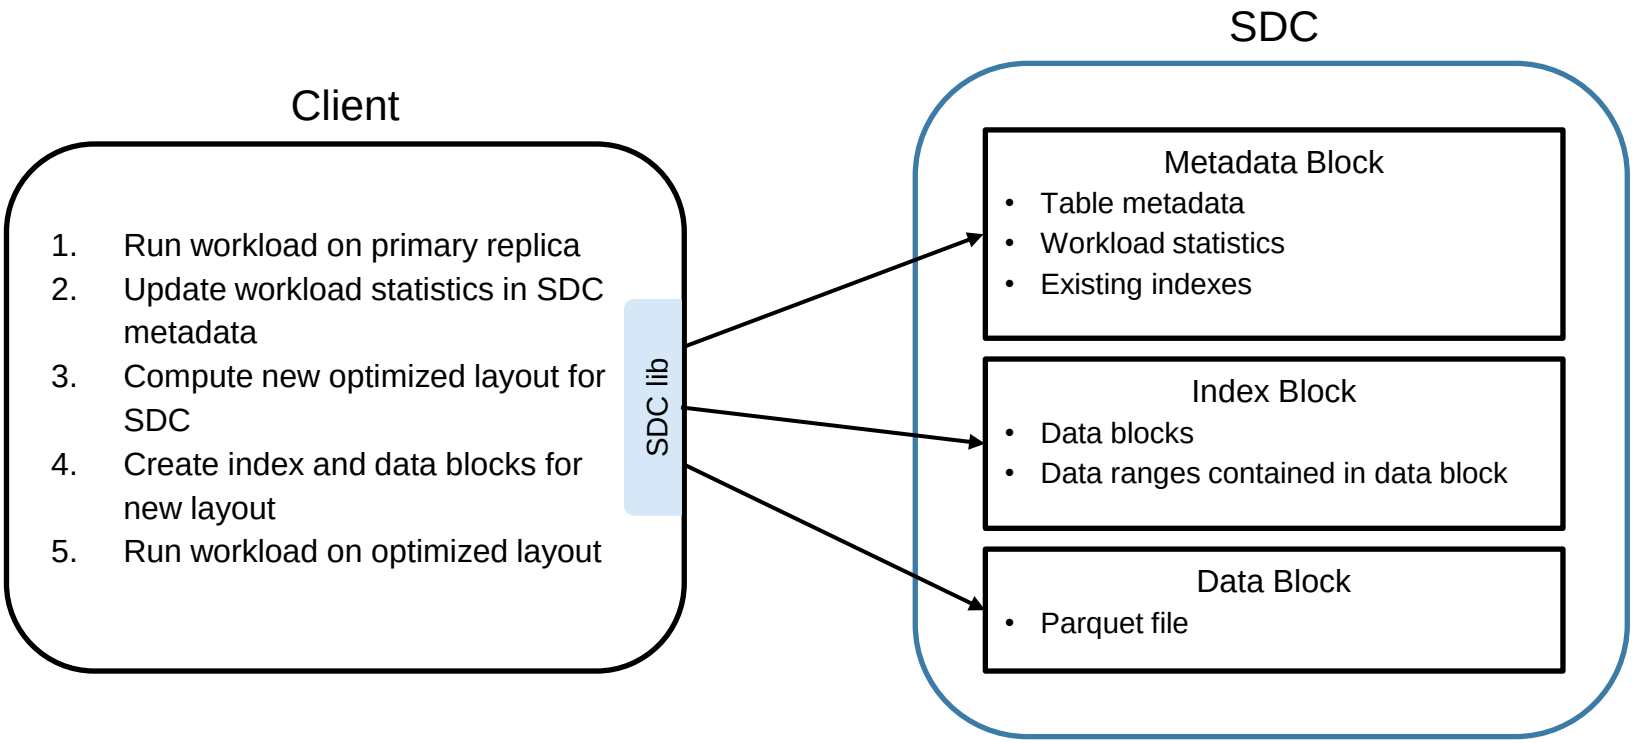

Demo of my SDC prototype

- SDC Library API: Projections and filters on tables
- Indexes: Primary, column range partitioning, Qd-Tree
- Dataset: NYC TLC Trip Record Data
- Workloads: Single-table range queries
- Storage layer: Local disk vs cloud storage (simulated)
- Data blocks: Apache Parquet files
- Metadata blocks: JSON files
- Clients: Single client

Benchmarks

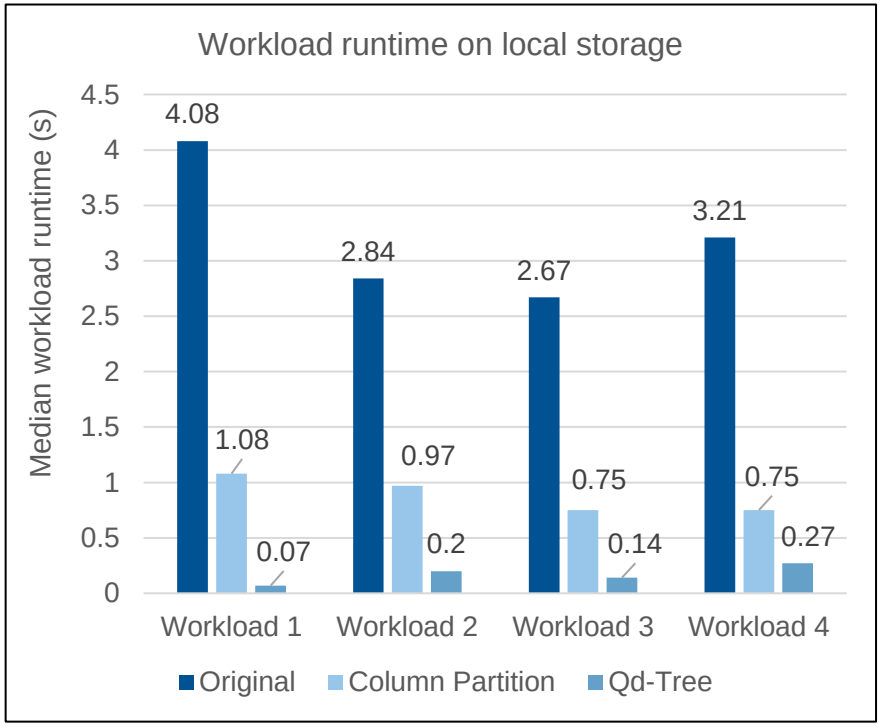
Workload runtime on local storage
Median workload runtime (s)
4.5
4
3.5
3
2.5
2
1.5
1
0.5
0
4.08
1.08
0.07
2.84
0.97
0.2
2.67
0.75
0.14
3.21
0.75
0.27
Workload 1
Workload 2
Workload 3
Workload 4
Original
Column Partition
Qd-Tree

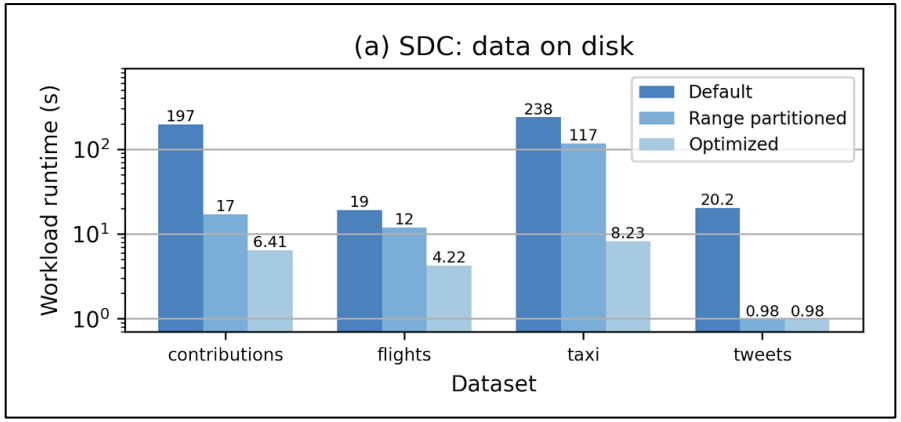
(a) SDC: data on disk
Workload runtime (s)
10^2
10^1
10^0
197
17
6.41
19
12
4.22
238
117
8.23
20.2
0.98
0.98
Default
Range partitioned
Optimized
contributions
flights
taxi
tweets
Dataset

Benchmarks

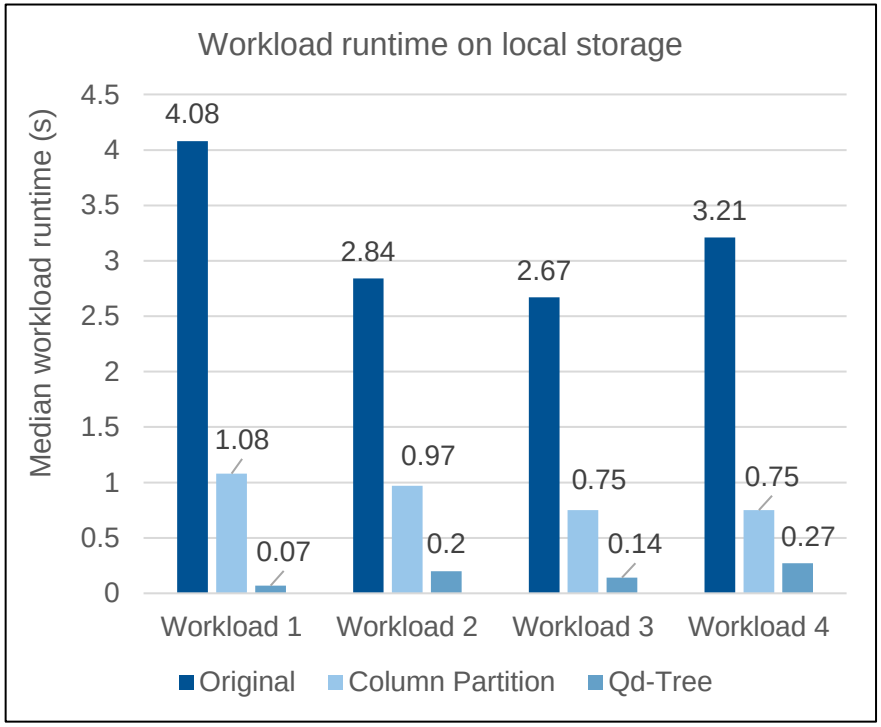
Workload runtime on local storage
Median workload runtime (s)
4.5
4
3.5
3
2.5
2
1.5
1
0.5
0
4.08
1.08
0.07
2.84
0.97
0.2
2.67
0.75
0.14
3.21
0.75
0.27
Workload 1
Workload 2
Workload 3
Workload 4
Original
Column Partition
Qd-Tree

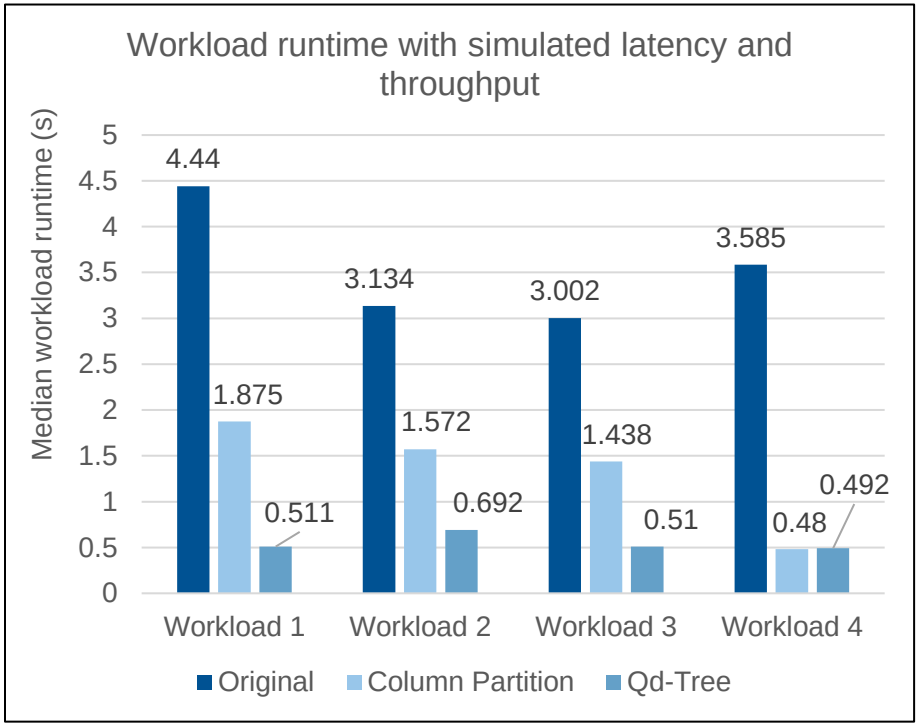
Workload runtime with simulated latency and throughput
Median workload runtime (s)
5
4.5
4
3.5
3
2.5
2
1.5
1
0.5
0
4.44
1.875
0.511
3.134
1.572
0.692
3.002
1.438
0.51
3.585
0.48
0.492
Workload 1
Workload 2
Workload 3
Workload 4
Original
Column Partition
Qd-Tree

Conclusion & further research

- Simple yet effective self-learned storage optimization
- Easy to add indexes
- Easy to integrate into any type of applications (not just DBMS)
- Use of data replication: trading off storage cost for query performance

Further research

- Distributing optimization work among clients
- Find query clusters in workload for effective indexes

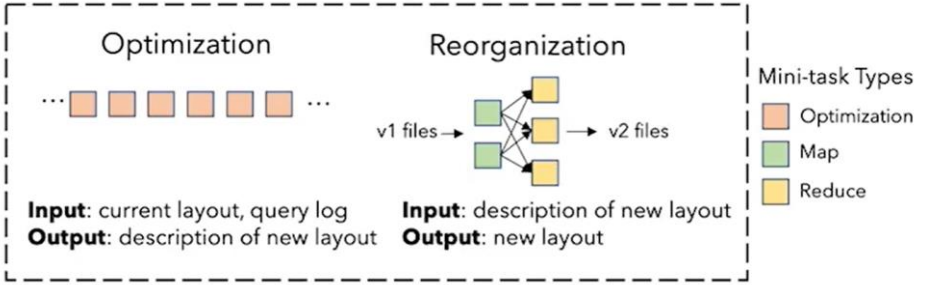

Code on Github:

https://github.com/thomasglas/SDCs

Benchmarking Qd-Tree min block sizes

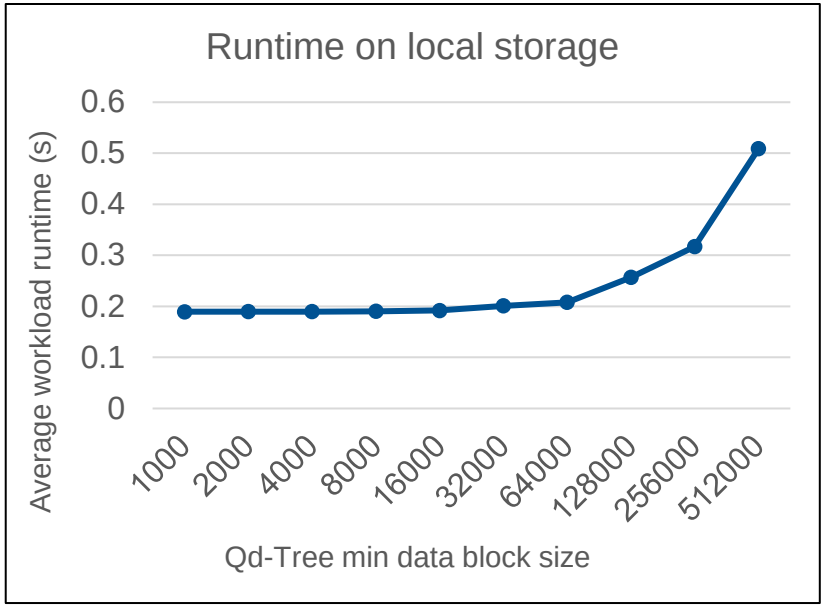

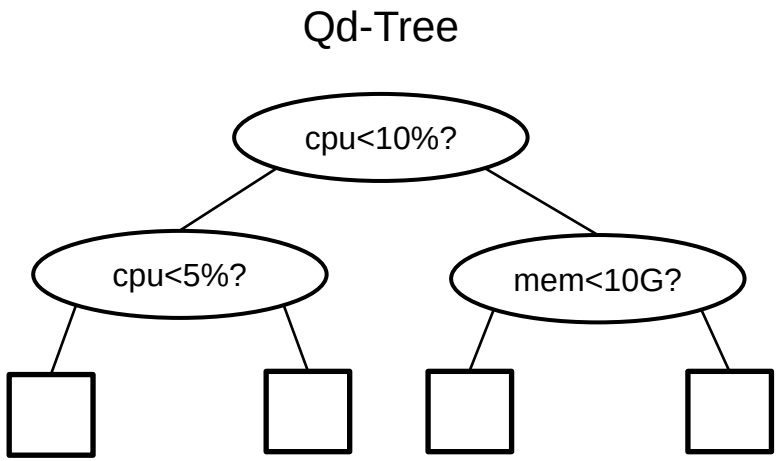

Benchmarking Qd-Tree min block sizes